

Problem__Saet__2

Yifan Li

Link to the GitHub

The link to my GitHub repository is <https://github.com/FYlee39/Stats-506/tree/main/PS2>.

Problem 1

a.

Version one:

```
#' Using for loop to implement the game
# '
#' @param n number of dice to roll
#' @return win total winnings
play_dice_v1 <- function(n){
  # cost 2 to play a roll
  win <- -2 * n
  roll_results <- sample(1: 6, n, replace=TRUE)
  for (single in roll_results){
    # if the dice shows 3 or 5, win the double points
    if (single == 3 | single == 5){
      win <- win + 2 * single
    }
  }
  return(win)
}

play_dice_v1(10)
```

```
[1] -4
```

Version two:

```
#' Using built-in R vectorized functions to implement the game
#'  
#' @param n number of dice to roll  
#' @return win total winnings  
play_dice_v2 <- function(n){  
  # cost 2 to play a roll  
  win <- -2 * n  
  roll_results <- sample(1: 6, n, replace=TRUE)  
  # For 3 and 5, let the winning be the double of itself  
  win_index <- (roll_results == 3) | (roll_results == 5)  
  roll_results[win_index] <- 2 * roll_results[win_index]  
  # For other numbers, set them to be 0  
  roll_results[!win_index] <- 0  
  win <- win + sum(roll_results)  
  return(win)  
}  
  
play_dice_v2(10)
```

```
[1] 20
```

Version three:

```
#' Using table to implement the game  
#'  
#' @param n number of dice to roll  
#' @return win total winnings  
play_dice_v3 <- function(n){  
  # cost 2 to play a roll  
  win <- -2 * n  
  roll_results <- table(sample(1: 6, n, replace=TRUE))  
  # number of 3 been rolled  
  num_three <- ifelse(!is.na(roll_results['3']), roll_results['3'], 0)  
  # number of 5 been rolled  
  num_five <- ifelse(!is.na(roll_results['5']), roll_results['5'], 0)  
  # total wining points  
  winning <- 3 * 2 * num_three + 5 * 2 * num_five
```

```

    win <- win + winning[[1]]
    return(win)
}

play_dice_v3(10)

```

[1] -10

Version four:

```

#' Using lapply to implement the game
#'
#' @param n number of dice to roll
#' @return win total winnings
play_dice_v4 <- function(n){
  # cost 2 to play a roll
  win <- -2 * n
  roll_results <- sample(1: 6, n, replace=TRUE)

  #' Get winning point of given roll
  #'
  #' @param x one roll result
  #' @return point the point gain from this rolling
  get_points <- function(x){
    point = 0
    # if the dice shows 3 or 5
    if (x == 3 | x == 5){
      # double the number to be the winning points
      point = x * 2
    }
    return(point)
  }

  winning <- sum(sapply(roll_results, get_points))
  win <- win + winning
  return(win)
}

play_dice_v4(10)

```

[1] -2

b.

Test for version one:

```
play_dice_v1(3)
```

```
[1] 0
```

```
play_dice_v1(3000)
```

```
[1] 1604
```

Test for version two:

```
play_dice_v2(3)
```

```
[1] 0
```

```
play_dice_v2(3000)
```

```
[1] 1954
```

Test for version three:

```
play_dice_v3(3)
```

```
[1] 0
```

```
play_dice_v3(3000)
```

```
[1] 2062
```

Test for version four:

```
play_dice_v4(3)
```

```
[1] 10
```

```
play_dice_v4(3000)
```

```
[1] 1532
```

c.

To demonstrate the same result, one needs set the same seed before each sampling. For 3 times of experiments:

```
set.seed(09152024)  
play_dice_v1(3)
```

```
[1] 4
```

```
set.seed(09152024)  
play_dice_v2(3)
```

```
[1] 4
```

```
set.seed(09152024)  
play_dice_v3(3)
```

```
[1] 4
```

```
set.seed(09152024)  
play_dice_v4(3)
```

```
[1] 4
```

For 3000 times of experiments:

```
set.seed(09152024)  
play_dice_v1(3000)
```

```
[1] 2344
```

```
set.seed(09152024)
play_dice_v2(3000)
```

[1] 2344

```
set.seed(09152024)
play_dice_v3(3000)
```

[1] 2344

```
set.seed(09152024)
play_dice_v4(3000)
```

[1] 2344

d.

For low input (1,000):

```
library(microbenchmark)
microbenchmark(v1=play_dice_v1(1000), v2=play_dice_v2(1000),
               v3=play_dice_v3(1000), v4=play_dice_v4(1000))
```

Unit: microseconds

expr	min	lq	mean	median	uq	max	neval	cld
v1	274.3	281.85	315.472	286.40	297.65	2636.7	100	a
v2	125.1	134.10	145.512	145.05	154.80	193.0	100	b
v3	351.5	381.45	400.431	400.55	416.05	484.2	100	c
v4	1259.2	1278.05	1343.738	1293.00	1344.20	3672.1	100	d

For large input (100,000):

```
library(microbenchmark)
microbenchmark(v1=play_dice_v1(100000), v2=play_dice_v2(100000),
               v3=play_dice_v3(100000), v4=play_dice_v4(100000))
```

Unit: milliseconds

expr	min	lq	mean	median	uq	max	neval	cld
v1	26.6637	30.90800	37.84597	33.99315	41.5251	166.5723	100	a
v2	10.8280	11.16655	11.94471	11.74655	12.1059	25.9205	100	b
v3	13.0132	13.47210	16.40812	13.93550	14.9918	154.6938	100	b
v4	132.2448	158.80760	177.16338	174.68360	187.4713	316.2325	100	c

From two experiments, one can find that among these four function, the implementation using built-in R vectorized functions is the fastest. Mean while, the function using `sapply` is the slowest.

e.

This game is unfair, to defend the decision using a Monte Carlo simulation, the version two will be used. There will be 100,000 times of experiments. Then the sample mean will be calculated, if the sample mean is no way near 0, then one can argue that this game is unfair.

```
sum <- 0
# Do 100,000 times of experiments, find the sample mean
n <- 100000
for (i in 1: n){
  sum <- sum + play_dice_v2(1)
}
sample_mean <- sum / n
sample_mean
```

```
[1] 0.6635
```

Since the sample mean is much greater than zero, one can argue that this is not a fair game.

Problem 2

a

```
raw_data <- read.csv('cars.csv', header=TRUE, sep=',',
                     col.names=c('Height',
                                   'Length',
                                   'Width',
```

```

'Driveline',
'Type',
'Hybird',
'Gears',
'Transmission',
'City_mpg',
'Fuel_type',
'Highway_mpg',
'Classification',
'ID',
'Make',
'Model_year',
'Year',
'horsepower',
'Torque'))
head(raw_data)

```

	Height	Length	Width	Driveline
1	140	143	202	All-wheel drive
2	140	143	202	Front-wheel drive
3	140	143	202	Front-wheel drive
4	140	143	202	All-wheel drive
5	140	143	202	All-wheel drive
6	91	17	62	All-wheel drive

	Type	Hybird	Gears
1	Audi 3.2L 6 cylinder 250hp 236ft-lbs	True	6
2	Audi 2.0L 4 cylinder 200 hp 207 ft-lbs Turbo	True	6
3	Audi 2.0L 4 cylinder 200 hp 207 ft-lbs Turbo	True	6
4	Audi 2.0L 4 cylinder 200 hp 207 ft-lbs Turbo	True	6
5	Audi 2.0L 4 cylinder 200 hp 207 ft-lbs Turbo	True	6
6	Audi 3.2L 6 cylinder 265hp 243 ft-lbs	True	6

	Transmission	City_mpg	Fuel_type	Highway_mpg
1	6 Speed Automatic Select Shift	18	Gasoline	25
2	6 Speed Automatic Select Shift	22	Gasoline	28
3	6 Speed Manual	21	Gasoline	30
4	6 Speed Automatic Select Shift	21	Gasoline	28
5	6 Speed Automatic Select Shift	21	Gasoline	28
6	6 Speed Manual	16	Gasoline	27

	Classification	ID	Make	Model_year	Year
1	Automatic transmission	2009	Audi A3 3.2	Audi 2009	Audi A3 2009
2	Automatic transmission	2009	Audi A3 2.0 T AT	Audi 2009	Audi A3 2009
3	Manual transmission	2009	Audi A3 2.0 T	Audi 2009	Audi A3 2009


```

4 Automatic transmission 2009 Audi A3 2.0 T Quattro Audi 2009 Audi A3 2009
5 Automatic transmission 2009 Audi A3 2.0 T Quattro Audi 2009 Audi A3 2009
6 Manual transmission 2009 Audi A5 3.2 Audi 2009 Audi A5 2009
  horsepower Torque
1      250      236
2      200      207
3      200      207
4      200      207
5      200      207
6      265      243

```

b.

```

gasoline_data <- raw_data[raw_data['Fuel_type'] == 'Gasoline', ]
head(gasoline_data)

```

```

  Height Length Width      Driveline
1    140    143   202 All-wheel drive
2    140    143   202 Front-wheel drive
3    140    143   202 Front-wheel drive
4    140    143   202 All-wheel drive
5    140    143   202 All-wheel drive
6     91     17    62 All-wheel drive

                                Type Hybird Gears
1              Audi 3.2L 6 cylinder 250hp 236ft-lbs  True    6
2 Audi 2.0L 4 cylinder 200 hp 207 ft-lbs Turbo  True    6
3 Audi 2.0L 4 cylinder 200 hp 207 ft-lbs Turbo  True    6
4 Audi 2.0L 4 cylinder 200 hp 207 ft-lbs Turbo  True    6
5 Audi 2.0L 4 cylinder 200 hp 207 ft-lbs Turbo  True    6
6              Audi 3.2L 6 cylinder 265hp 243 ft-lbs  True    6

      Transmission City_mpg Fuel_type Highway_mpg
1 6 Speed Automatic Select Shift      18 Gasoline      25
2 6 Speed Automatic Select Shift      22 Gasoline      28
3              6 Speed Manual      21 Gasoline      30
4 6 Speed Automatic Select Shift      21 Gasoline      28
5 6 Speed Automatic Select Shift      21 Gasoline      28
6              6 Speed Manual      16 Gasoline      27

      Classification                                ID Make  Model_year Year
1 Automatic transmission                2009 Audi A3 3.2 Audi 2009 Audi A3 2009
2 Automatic transmission                2009 Audi A3 2.0 T AT Audi 2009 Audi A3 2009

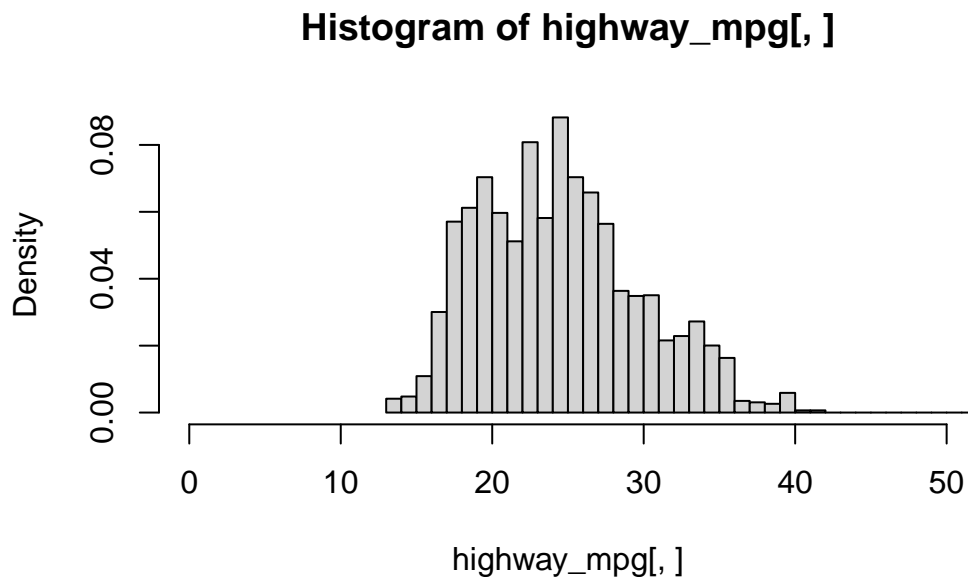
```

3	Manual transmission	2009 Audi A3 2.0 T Audi	2009 Audi A3 2009
4	Automatic transmission	2009 Audi A3 2.0 T Quattro Audi	2009 Audi A3 2009
5	Automatic transmission	2009 Audi A3 2.0 T Quattro Audi	2009 Audi A3 2009
6	Manual transmission	2009 Audi A5 3.2 Audi	2009 Audi A5 2009
	horsepower	Torque	
1	250	236	
2	200	207	
3	200	207	
4	200	207	
5	200	207	
6	265	243	

c.


The original data distribution is:

```
highway_mpg <- gasoline_data['Highway_mpg']
hist(highway_mpg[,], breaks = 200, probability = TRUE, xlim = c(0, 50))
```



Since the data are all positive and the distribution has a right skew with a long tail, a log transformation would likely be the best choice. Then update the data in the data frame.

```
transformed_highway_mpg <- log10(highway_mpg)
hist(transformed_highway_mpg[,], breaks = 20, probability = TRUE, xlim = c(1, 2.5))
```



```
gasoline_data['Highway_mpg'] <- transformed_highway_mpg
```

d.

```
# Make Year to be categorical variable
gasoline_data$Year <- as.factor(gasoline_data$Year)
d_lm <- lm(Highway_mpg~Torque + horsepower
          + Height + Length + Width
          + Year, data=gasoline_data)
summary(d_lm)
```

Call:

```
lm(formula = Highway_mpg ~ Torque + horsepower + Height + Length +
    Width + Year, data = gasoline_data)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.23782	-0.04076	-0.00180	0.04297	1.05035

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	1.523e+00	9.625e-03	158.236	< 2e-16 ***
Torque	-9.964e-04	2.934e-05	-33.956	< 2e-16 ***
horsepower	4.012e-04	3.033e-05	13.227	< 2e-16 ***
Height	1.759e-04	1.501e-05	11.719	< 2e-16 ***
Length	1.509e-05	1.177e-05	1.282	0.19980
Width	-3.788e-05	1.205e-05	-3.144	0.00168 **
Year2010	-9.473e-03	9.015e-03	-1.051	0.29342
Year2011	-1.055e-03	9.000e-03	-0.117	0.90665
Year2012	1.742e-02	9.071e-03	1.921	0.05485 .

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.0613 on 4582 degrees of freedom

Multiple R-squared: 0.5638, Adjusted R-squared: 0.563

F-statistic: 740.3 on 8 and 4582 DF, p-value: < 2.2e-16

From the coefficient of torque, which is -9.964155×10^{-4} , meaning for each additional unit of torque, highway MPG will decrease by -9.964155×10^{-4} while holding other variables constant. The coefficient is significant with a p-value less than 2×10^{-16} , indicating that the relationship is statistically significant.

e.

```
e_lm <- lm(Highway_mpg ~ horsepower * Torque +
           + Height + Length + Width, data=gasoline_data)
summary(e_lm)
```

Call:

```
lm(formula = Highway_mpg ~ horsepower * Torque + +Height + Length +
    Width, data = gasoline_data)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.23084	-0.03530	-0.00158	0.03424	1.06415

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	1.673e+00	6.746e-03	247.952	< 2e-16	***
horsepower	-7.294e-05	3.359e-05	-2.172	0.0299	*
Torque	-1.562e-03	3.352e-05	-46.618	< 2e-16	***
Height	1.273e-04	1.414e-05	9.002	< 2e-16	***
Length	1.042e-05	1.102e-05	0.946	0.3443	
Width	-4.992e-05	1.129e-05	-4.422	9.98e-06	***
horsepower:Torque	1.710e-06	6.139e-08	27.864	< 2e-16	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1


Residual standard error: 0.05747 on 4584 degrees of freedom
Multiple R-squared: 0.6165, Adjusted R-squared: 0.616
F-statistic: 1228 on 6 and 4584 DF, p-value: < 2.2e-16

For the three different horsepower values, they will be the min, mean and the max value.

```
library(interactions)

horsepower_range <- c(min(gasoline_data$horsepower),
                      mean(gasoline_data$horsepower),
                      max(gasoline_data$horsepower))

interact_plot(e_lm, pred=Torque, modx=horsepower, data=gasoline_data,
              modx.values=horsepower_range,
              main.title="Interaction Between Torque and Horsepower")
```



f.

```
# create the design matrix
design_matrix <- model.matrix(Highway_mpg ~ Torque + horsepower
                             + Height + Length + Width
                             + Year, data=gasoline_data)

Y <- gasoline_data$Highway_mpg

design_matrix_t <- t(design_matrix)

XTX_inv <- solve(design_matrix_t %*% design_matrix)

XTY <- design_matrix_t %*% Y

beta_hat <- XTX_inv %*% XTY
# Compare two values
beta_hat
```

```
      [,1]
(Intercept) 1.523037e+00
```

Torque	-9.964155e-04
horsepower	4.012067e-04
Height	1.758848e-04
Length	1.509263e-05
Width	-3.788045e-05
Year2010	-9.473036e-03
Year2011	-1.055492e-03
Year2012	1.742184e-02

```
manual_coef <- setNames(as.vector(beta_hat), names(d_lm$coefficients))
all.equal(manual_coef, d_lm$coefficients)
```

```
[1] TRUE
```

The result is True, which shows that the manual result is the same with the `lm` result.